



Dell Red Hat Cloud Solutions Hardware Deployment Guide

A Dell Deployment Guide

Dell Revolutionary Cloud and Big Data Solutions Engineering
October 27, 2014

Revisions

Date	Description
10/27/2014	Initial release version.

This document is for informational purposes only and may contain typographical errors and technical inaccuracies. The content is provided as is, without express or implied warranties of any kind.



Table of Contents

- Revisions 2
- Table of Contents 3
- List of Tables 4
- Trademarks 5
- Notes, Cautions, and Warnings 6
- Overview..... 7
 - Intended Audience 7
- Configuring R620, R720 and R720xd Hardware 8
 - IPMI Configuration 8
 - User information 8
 - Network configuration..... 8
 - IPMI Validation..... 8
 - BIOS Configuration 9
 - Infrastructure and Virtualization Node Settings 9
 - Storage Node Settings 9
 - BIOS Settings for all Nodes 10
 - RAID Configuration 11
 - Dell R620 and R720xd RAID and BIOS Configuration Overview 11
- EqualLogic Storage Group 12
- Network Configuration 13
 - Configuring the Dell Force 10 S55 and/or S4810 13
- References 16
- To Learn More 16

List of Tables

Table 1: R620 Virtualization Settings.....9

Table 2: R720xd Storage Settings.....9

Table 3: BIOS Settings10

Table 4: EqualLogic Information Needed from Configuration12



Trademarks

© 2014 Dell Inc. All rights reserved. Reproduction of this material in any manner whatsoever without the express written permission of Dell Inc. is prohibited. For more information, contact Dell. Dell, the Dell logo, and the Dell badge, and PowerEdge are trademarks of Dell Inc.

Trademarks used in this text: Dell™, the DELL logo, Dell Precision™, OptiPlex™, Latitude™, PowerEdge™, PowerVault™, PowerConnect™, OpenManage™, EqualLogic™, KACE™, FlexAddress™, Force10™ and Vostro™ are trademarks of Dell Inc. Intel®, Pentium®, Xeon®, Core™ and Celeron® are registered trademarks of Intel Corporation in the U.S. and other countries. AMD® is a registered trademark and AMD Opteron™, AMD Phenom™, and AMD Sempron™ are trademarks of Advanced Micro Devices, Inc. Microsoft®, Windows®, Windows Server®, MS-DOS® and Windows Vista® are either trademarks or registered trademarks of Microsoft Corporation in the United States and/or other countries.

Red Hat®, Red Hat® Enterprise Linux®, the Shadowman logo, and JBoss are trademarks of Red Hat, Inc., registered in the U.S. and other countries. Linux® is the registered trademark of Linus Torvalds in the U.S. and other countries. Novell® is a registered trademark and **SUSE™ is a trademark of Novell Inc. in the United States and other countries.** Oracle® is a registered trademark of Oracle Corporation and/or its affiliates. Citrix®, Xen®, XenServer® and XenMotion® are either registered trademarks or trademarks of Citrix Systems, Inc. in the United States and/or other countries. VMware®, Virtual SMP®, vMotion®, vCenter®, and vSphere® are registered trademarks or trademarks of VMware, Inc. in the United States or other countries.

DISCLAIMER: The OpenStack® Word Mark and OpenStack Logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries, and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation or the OpenStack community.

Other trademarks and trade names may be used in this publication to refer to either the entities claiming the marks and names or their products. Dell Inc. disclaims any proprietary interest in trademarks and trade names other than its own.



Notes, Cautions, and Warnings



A **Note** indicates important information that helps you make better use of your system.



A **Caution** indicates potential damage to hardware or loss of data if instructions are not followed



A **Warning** indicates a potential for property damage, personal injury, or death.

This document is for informational purposes only and may contain typographical errors and technical inaccuracies. The content is provided as is, without express or implied warranties of any kind.



Overview

This guide provides information necessary to deploy the Dell Red Hat Cloud Solutions with Red Hat Enterprise Linux™ OpenStack Platform 5, on Dell R620, R720, and R720xd servers with Dell PowerEdge H710 disk controller; and the network with the Dell Force10 S4810 and the S55 Switches.

Intended Audience

This guide assumes the reader is familiar with:

- OpenStack
- PowerEdge R620 and R720xd RAID and BIOS configuration
- Red Hat Enterprise Linux (RHEL)
- Network Configuration



Configuring R620, R720 and R720xd Hardware

This section describes manually configuring Dell R720 and the Dell R720xd server hardware for the Dell Red Hat Cloud Solutions with Red Hat Enterprise Linux™ OpenStack Platform:

- IPMI Configuration (see page 8)
- BIOS Configuration (see page 9)
- RAID Configuration (see page 11)

IPMI Configuration

You must perform configuration of the iDRAC on supported systems. Dell recommends that you configure these settings from the iDRAC web interface, or directly on the node console.

User information

1. Set credentials for the *root* user including changing the password to a new one based on good password standards.
2. Set privileges for the user to *Admin* level privileges including over LAN
3. *Enable* the user, if disabled.

Network configuration

1. Sets the iDRAC IP address, subnet mask, default gateway IP, and default VLAN (id is *10*).
2. Sets the iDRAC IP address source to be *static* IP addressing.
3. Sets the iDRAC NIC mode to *Dedicated* from *Shared*.

IPMI Validation

As part of the IPMI setup, validation that remote commands can be executed is essential.

1. Install to your workstation the IPMI Utilities from SourceForge:
 - a. For Linux - <http://sourceforge.net/projects/ipmitool/>
 - b. For Windows - <http://ipmiutil.sourceforge.net/>
 - c. Validate that you have all the requirements, and that it will run.
2. Plug your Ethernet port into a switch port that has the same vLAN as your iDRACs.
3. Configure your network to use an IP address in the iDRAC network range.
4. Configure the IPMI over LAN Setting to *Enabled*.
5. Execute an IPMI command:
 - a. `for i in $(seq 162 170); do ipmitool -P "password" -U "username" -H 192.168.200.$i power status; done`
This will perform a simple, non-destructive poll of the iDRAC from *192.168.200.162* to *192.168.200.170*, and then return the power status of *on* or *off*.
 - b. You can replace the keyword status with *reset*, *off*, or *on*.
6. Ensure that all machines return responses to the command.



BIOS Configuration

The primary BIOS configurations for the Dell Red Hat cloud solutions with RHEL OpenStack Platform are for Virtualization/Infrastructure nodes and Storage nodes.

For more information about Dell PowerEdge R720xd BIOS settings, please see the [Dell PowerEdge R720 and R720xd Owner's Manual](#).

Infrastructure and Virtualization Node Settings

This section describes the settings for nodes that will be used on the Solution Admin Host, OpenStack controllers and Compute nodes. For the R620 the settings below need to be double-checked.

Table 1: R620 Virtualization Settings

Type	Setting	State
Integrated Devices	SR-IOV Global Enable	Enabled
Processor Setting	Execute Disable	Enabled
Processor Setting	Virtualization Technology	Enabled
Boot Settings	Boot Mode	BIOS
Boot Setting	Boot Order	BIOS Boot mode Set PXE boot first local drive second

Storage Node Settings

This section describes storage settings for the PowerEdge R720xd platform used as a storage node. You must ensure these settings are correct.

Table 2: R720xd Storage Settings

Type	Setting	State
Integrated Devices	SR-IOV Global Enable	Disabled
Processor Setting	Execute Disable	Enabled
Processor Setting	Virtualization Technology	Disabled
Boot Settings	Boot Mode	BIOS
BIOS	Boot Order	BIOS Boot mode Set PXE boot first local drive second



BIOS Settings for all Nodes

This section outlines all the settings that Dell recommends. The quickest way to put these in place is to press "F3", and then validate them.

Table 3: BIOS Settings

BIOS Category	Setting	State
Memory Settings	System Memory Testing	Disabled
Memory Settings	Memory Operating Mode	Optimizer Mode
Memory Settings	Node Interleaving	Disabled
Processor Settings	Logical Processor	Enabled
Processor Settings	QPI Speed	Maximum Data Rate
Processor Settings	Alternate RTID Setting	Disabled
Processor Settings	Adjacent Cache Line Prefetch	Enabled
Processor Settings	Hardware Prefetcher	Enabled
Processor Settings	DCU Streamer Prefetcher	Enabled
Processor Settings	DCU IP Prefetcher	Enabled
Processor Settings	Logical Processor Idling	Disabled
Processor Settings	Dell Controlled Turbo	Disabled
Processor Settings	Number of cores per Processor	All
Processor Settings	Processor 64-bit Support	Yes
Integrated Devices	Integrated RAID Controller	Enabled
Integrated Devices	OS Watchdog Timer	Disabled
Integrated Devices	I/OAT DMA Engine	Enabled
Integrated Devices	Memory Mapped I/O above 4GB	Enabled
System Profile Settings	System Profile	Performance



RAID Configuration

You must perform RAID configuration manually, using the option ROM for the RAID controller. For more information on configuring your controller please see the [Dell PowerEdge RAID Controller \(PERC\) H310, H710, H710P, and H810 User's Guide](#).

- Virtualization/Infrastructure nodes are RAID 10, spanning all drives on the system.
- Storage Nodes utilize:
 - The two back-side disks are to be set as a RAID-1 pair
 - Add the front-side disks, SSD and HDD, are to be setup a single disk RAID 0. In other words, a RAID 0 virtual disk to be created for each of the front-side disks.

Dell R620 and R720xd RAID and BIOS Configuration Overview

You can use the Dell Lifecycle controller (http://ftp.dell.com/Manuals/all-products/esuprt_electronics/esuprt_software/esuprt_remote_ent_sys_mgmt/esuprt_rmte_ent_sys_lifecycle_cntrlr/) to configure all the above functions. This is a brief overview of the steps.

1. Boot up each machine.
2. At the boot screen Press [*F10*], then *Enter Lifecycle Controller*.
3. Select *System Setup*.
4. Select *Advanced Hardware Configuration*.
5. Select *System BIOS*.
6. Make the changes per the charts above.
7. Press [*Esc*] to exit back.
8. Save your changes when prompted.
9. Select *iDRAC*.
10. Configure iDRAC Network settings.
11. Press [*Esc*] to exit back.
12. Save your changes when prompted.
13. Select *Configure Devices*.
14. Select *RAID Controller*.
15. Configure your Virtual Disks.
16. Press [*Esc*].
17. Save your changes when prompted. You may see a repeated message that asks you to *press Ctrl-Alt-Del* in order to reboot.
18. Reboot.



EqualLogic Storage Group

The EqualLogic Storage Group can consist of one or more storage arrays with one or more storage groups. The configuration of the arrays is beyond the scope of this document. Please refer to [EqualLogic Support Website \(https://eqlsupport.dell.com/secure/login.aspx\)](https://eqlsupport.dell.com/secure/login.aspx) for the latest guides, whitepapers and best practices on how to setup your Storage Group for your application.

Once the Storage Group(s) are setup, the information contained in Table 4 must be collected to configure your software. The `san_thin_provision` variable should be left at the default if available; if not, then it must be set to `False`. Refer to https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux_OpenStack_Platform/5/html/Configuration_Reference_Guide/section_volume-drivers.html#dell-equallogic-driver for the current list.

Table 4: EqualLogic Information Needed from Configuration

[DEFAULT]	Description
<code>volume_driver = cinder.volume.drivers.eqlx</code>	Dell EqualLogic volume driver
<code>san_ip = <IP_address_of_EQLX></code>	IP address used to reach the EqualLogic Group through SSH
<code>san_login = <user_name></code>	User name to login to the Group manager via SSH at the <code>san_ip</code>
<code>san_password = <password></code>	Password to login to the Group manager via SSH at the <code>san_ip</code> (not used when <code>san_private_key</code> is set)
<code>san_thin_provision = <true false></code>	Enable/disable creation of thin-provisioned volumes
<code>san_ssh_port = 22</code>	Port used for SSH
<code>ssh_conn_timeout = 30</code>	Timeout value, in seconds, used by CLI commands over SSH
<code>san_private_key = <filename></code>	Filename of the private key used for SSH authentication
<code>ssh_min_pool_conn = 1</code>	Minimum number of SSH connections in the pool
<code>ssh_max_pool_conn = 5</code>	Maximum number of SSH connections in the pool
<code>eqlx_chap_login = admin</code>	Existing CHAP account name
<code>eqlx_chap_password = password</code>	Password for specified CHAP account name
<code>eqlx_cli_max_retries = 5</code>	Maximum retry count for reconnection
<code>eqlx_cli_timeout = 30</code>	Timeout for the Group Manager cli command execution
<code>eqlx_group_name = group-0</code>	Group name to use for creating volumes
<code>eqlx_pool = default</code>	Pool in which volumes will be created
<code>eqlx_use_chap = False</code>	Use CHAP authentication for targets?



Network Configuration

If you are not using Dell Force 10 S55 and/or S4810 switches, you will need to program the switches to support your cloud instantiation. Your switches are expected to support the following:

- Support for IEEE 802.1Q vLAN traffic and port tagging
- Support using one untagged and multiple tagged vLANs on the same port
- Ability to provide a minimum of 170 Gigabit Ethernet ports in a non-blocking configuration within Provisioning vLAN
 - Configuration can be a single switch or a combination of stacked switches to meet the additional requirements
- The ability to create link aggregation groups (LAGs) with a minimum of two physical links in each LAG
- If multiple switches are stacked:
 - The ability to create a LAG across stacked switches
 - Full-bisection bandwidth
 - Support for vLANs to be available across all switches in the stack
- 250,000 packets-per-second capability per switch
- A managed switch that supports SSH and serial line configuration
- SNMP v3 support

Configuring the Dell Force 10 S55 and/or S4810

The bundle workbooks are designed to help configure the cluster. Follow the steps outlined to setup your network.

1. Fill each workbook page with all information marked required:
 - a. On the switch config page fill in each device name, NIC connected to port
 - b. Using your planning sheet and vLAN information, mark which vLAN is either untagged or tagged for each port connected. Not all vLAN will be used on each port.
2. Using the workbook you can now create your configuration:
 - a. Individual Port Config Example for S55

```
no ip address
portmode hybrid
switchport
spanning-tree rstp edge-port
```

- b. Individual Port Config example for S4810

```
no ip address
portmode hybrid
switchport
flowcontrol rx on tx off
spanning-tree rstp edge-port
```



- c. Configuring VLANs. With option *portmode hybrid* on each port, you can program one VLAN to be untagged and other VLANs as tagged. This will send all untagged traffic to the VLAN that is marked untagged. For example, DHCP Request will be placed on that VLAN, and a DHCP server listening on that VLAN will be able to respond.

Example VLAN config with overlapping tagged and untagged ports:

```
Int vlan 100
no ip address
! From your Spread Sheet
tagged TenGigabitEthernet 0/1 - 47
untagged TenGigabitEthernet 0/0
!
int vlan 200
tagged TenGigabitEthernet 0/0-47
int vlan 300
untagged TenGigabitEther 0/1-10
!
!
int vlan 300
tagged TenGigabitEthernet 0/0-47
int vlan 300
untagged TenGigabitEther 0/1-10
!
```

- d. General switch options for the S55 and S4810:

```
!
hostname YOUR_HOSTNAME
!
username Tome password 7 MYpassword 1 privilege 15
!
protocol spanning-tree rstp
no disable
hello-time 1
max-age 6
forward-delay 4
bridge-priority 16384
!
no ip telnet server enable
!
arp learn-enable
!
clock timezone EST -5
!
ip ssh server enable
!
protocol lldp
advertise dot1-tlv port-protocol-vlan-id
advertise dot3-tlv max-frame-size
advertise management-tlv system-description system-name
advertise med
```



```
no disable
!  
reload-type normal-reload  
!
```

e. How to create a VLT Domain on the S4810:

```
!  
vlt domain 899  
peer-link port-channel 100  
primary-priority 4096  
unit-id 0  
!  
!  
interface Port-channel 100  
description Connection to VLT peer  
no ip address  
channel-member fortyGigE 0/48  
no shutdown  
!
```



You do not need to add the port-channel or the vlt to a vLAN. All vLANs will be passed on it.

References

Additional information can be obtained at www.dell.com/openstack or by e-mailing openstack@dell.com.

If you need additional services or implementation help, please contact your Dell sales representative.

To Learn More

For more information on the Dell Red Hat Cloud Solutions with Red Hat Enterprise Linux™ OpenStack Platform visit:

<http://www.dell.com/openstack>

©2014 Dell Inc. All rights reserved. Trademarks and trade names may be used in this document to refer to either the entities claiming the marks and names or their products. Specifications are correct at date of publication but are subject to availability or change without notice at any time. Dell and its affiliates cannot be responsible for errors or omissions in typography or photography. **Dell's Terms and Conditions of Sales and Service apply and are available on request. Dell service offerings do not affect consumer's statutory rights.**

Dell, the DELL logo, and the DELL badge, PowerConnect, and PowerVault are trademarks of Dell Inc.

